

Modern Hopfield Networks: A Short Note

Claude (Anthropic) and Khanh Nguyen

May 2026

1 Summary

Modern Hopfield Networks (MHN) are classical Hopfield networks with the energy function swapped from quadratic to exponential. Two things change as a result: storage capacity goes from linear in dimension to exponential in dimension, and the update rule becomes literally the attention operation used in transformers.

2 Classical Hopfield (1982)

N binary neurons $s \in \{-1, +1\}^N$. Store M patterns $\xi^1, \dots, \xi^M \in \{-1, +1\}^N$ via the Hebbian outer-product rule:

$$W = \frac{1}{N} \sum_{\mu=1}^M \xi^\mu (\xi^\mu)^\top.$$

Energy and dynamics:

$$E(s) = -\frac{1}{2} s^\top W s, \quad s_i \leftarrow \text{sign}((Ws)_i).$$

Each update flips one neuron only if it decreases energy, so dynamics converge to a local minimum. The stored patterns ξ^μ are minima provided the patterns are sufficiently distinct.

Capacity. Amit, Gutfreund, and Sompolinsky (1985) showed via replica methods that the network can reliably store

$$M \approx 0.138 N$$

patterns. Above this, cross-pattern interference (spurious minima created by pattern superposition in W) destroys retrieval.

Why capacity is limited. All patterns are summed into the single matrix W . Every pair of patterns contributes a cross-term $\xi^\mu (\xi^\nu)^\top$. As M grows, these cross-terms accumulate and the energy landscape develops spurious minima between true patterns. The quadratic energy is too coarse to keep wells separated.

3 Modern Hopfield (Demircigil et al. 2017; Ramsauer et al. 2020)

Two changes:

1. States are continuous: $s \in \mathbb{R}^d$ rather than $\{-1, +1\}^N$.
2. Energy is **exponential** (technically a log-sum-exp) rather than quadratic.

With patterns stacked as columns of $\Xi = [\xi^1, \dots, \xi^M] \in \mathbb{R}^{d \times M}$:

$$E(s) = -\frac{1}{\beta} \log \sum_{\mu=1}^M \exp(\beta \xi^\mu \cdot s) + \frac{1}{2} \|s\|^2 + \text{const.}$$

The negative log-sum-exp is a smooth *maximum*. It carves a sharp well around each stored pattern ξ^μ , with sharpness controlled by the inverse-temperature parameter β . The $\|s\|^2$ term keeps the state bounded.

3.1 Update rule

Minimising $E(s)$ from a query s via the Concave–Convex Procedure gives one update step:

$$s_{\text{new}} = \Xi \cdot \text{softmax}(\beta \Xi^\top s).$$

That is: compute affinities $\beta \langle s, \xi^\mu \rangle$, take the softmax to get retrieval weights p_μ , return the weighted sum of patterns $\sum_\mu p_\mu \xi^\mu$.

3.2 Connection to attention

If we separate “stored cues” from “retrieved content” (let Ξ be the keys K and let the values V be different from K), the update rule becomes:

$$\text{Attn}(q, K, V) = V \cdot \text{softmax}(\beta K^\top q), \quad \beta = \frac{1}{\sqrt{d}}.$$

One attention layer in a transformer is exactly one update step on a modern Hopfield energy. Multi-head attention is parallel retrieval from multiple associative memories. The $\sqrt{d}$ scaling that transformer papers introduce as “numerical stabilisation” is, in this framing, the natural inverse-temperature scaling needed to keep the Boltzmann distribution in the non-trivial regime.

3.3 Why capacity is exponential

The classical Hopfield’s quadratic energy summed pattern overlaps *additively*: pattern μ ’s well bled into pattern ν ’s as soon as $\xi^\mu \cdot \xi^\nu$ was non-negligible. The exponential energy makes wells *multiplicatively* sharp: small angular separation between patterns produces vastly different exp values, so wells stay disjoint until patterns are very close.

For random patterns drawn uniformly on the unit sphere in $\mathbb{R}^d$, near-orthogonality holds for $\sim \exp(\Theta(d))$ vectors (Johnson–Lindenstrauss / sphere packing). The log-sum-exp energy supports that many distinct wells. Demircigil et al. give the formal bound:

$$M \lesssim \exp(d/4)$$

patterns can be stored with high probability of correct retrieval.

Concrete numbers for transformers. Head dimension d in modern LLMs is typically 64–256:

d	approximate capacity $\sim \exp(d/4)$
64	$\sim 10^7$
128	$\sim 10^{14}$
256	$\sim 10^{28}$

At $d = 128$ the capacity is $\sim 10^{14}$, far exceeding any context length a transformer would actually run. **For practical purposes, attention has unlimited memory capacity per head.** Whatever caps transformer context length, it is not the attention operation’s storage capacity — it is compute (the $O(N^2)$ attention math), memory bandwidth, or training-data distribution.

4 What retrieval looks like geometrically

Given a query q and stored patterns $\{\xi^\mu\}$:

- If q is near one stored ξ^* and far from the rest, the softmax saturates to a one-hot on ξ^* and retrieval returns ξ^* exactly.
- If q is between two patterns of comparable affinity, the softmax blends them and retrieval returns a weighted average.
- If q is far from everything, the softmax approaches uniform and retrieval returns the centroid of all patterns (uninformative).

The $\sqrt{d}$ temperature keeps the operation in a “sharp but graceful” regime: clean retrieval on a single match, smooth blending on near-ties, soft fallback when nothing matches. Hard argmax (zero temperature) would be brittle; uniform mixing (infinite temperature) would be useless. Softmax with $\beta = 1/\sqrt{d}$ is the unique well-behaved middle.

5 One step vs. iterative retrieval

Modern Hopfield retrieval typically **converges in one update step** for well-separated patterns. This is why attention works as a single-layer operation in transformers — no iteration required.

For ambiguous queries (multiple stored patterns about equally close), additional update steps would refine the answer. Transformers do not iterate retrieval; instead they stack many attention layers, each performing one update step on a (different) associative store learned during training.

6 What this framing buys you

1. **Capacity bounds for transformers.** Per-head storage is $\sim \exp(d_{\text{head}}/4)$. For modern head dimensions this is effectively unbounded.
2. **Predicting behaviour.** When does attention “look up” a single key versus “blend” multiple? It depends on the angular geometry of keys around the query — analyzable via Hopfield tools.
3. **New architectures.** Hopfield layers can be used as standalone associative-memory modules, generalising attention to non-sequence contexts (graph nets, set inputs, set-to-set mappings).
4. **Theoretical bridge.** Connects transformers to the 40-year tradition of statistical-mechanics analysis of neural networks (spin glasses, replica method, mean-field theory).

7 What this framing does not buy you

- It does not explain *why* transformers learn good keys/queries — it analyses retrieval *given* the keys and queries.
- It does not explain depth — why do many stacked layers do something irreducible?
- It does not explain attention-head specialisation — that is an empirical/training story.
- It does not directly address in-context learning, scaling laws, or instruction tuning — those need separate frameworks.

The modern-Hopfield framing is the right abstraction for the **single-layer mechanics** of attention. Multi-layer dynamics, in-context learning, and the broader question of why transformers work require additional theory. But for the specific question “what is the attention operation, mathematically?” — it is content-addressable memory retrieval on an exponential-capacity associative store, and that statement is exact.

8 Is the Hopfield Nobel “really physics”?

The 2024 Nobel Prize in Physics went to John Hopfield and Geoffrey Hinton “for foundational discoveries and inventions that enable machine learning with artificial neural networks.” A number of physicists publicly grumbled that this was not physics. The grumbling is sociologically understandable (these people did not appear at APS conferences) but mathematically off-base.

- Hopfield holds a physics PhD from Cornell, worked at Bell Labs in solid-state physics, and spent decades as a physics professor at Princeton. His 1982 paper applies Ising-model formalism and statistical-mechanics energy methods directly.
- The capacity results (Amit, Gutfreund, Sompolinsky 1985) were derived using *replica methods* from spin-glass theory — the same toolkit Parisi used for his 2021 Nobel.
- Hinton’s Boltzmann machines (which contributed to his share of the 2024 Nobel) are named after Boltzmann distributions and use partition functions, free energies, and Monte Carlo sampling. Not borrowed as metaphor; borrowed as machinery.

The math is statistical mechanics. The constructions are energy-function-driven dynamical systems. The capacity bounds are proven with techniques developed for disordered systems. The complaint that “this isn’t real physics” is about the people, not the content.

9 Hopfield capacity vs. the information-theoretic bound

A clean way to see what modern Hopfield achieves: compare its capacity to the Shannon-style information bound.

A d -dimensional binary-precise state can index at most 2^d distinct patterns:

$$M_{\text{Shannon}} \leq 2^d \quad \Leftrightarrow \quad \log_2 M \leq d.$$

Modern Hopfield with random patterns achieves:

$$M_{\text{MHN}} \lesssim e^{d/4} = 2^{d/(4 \ln 2)} \approx 2^{0.36 d}.$$

So MHN packs about **36% of the bits-per-dimension of the information-theoretic ceiling**. Not optimal — but within a constant factor in the exponent.

Why this specific ratio. This is exactly what Shannon’s *random coding bound* predicts. For random codewords on the d -sphere, the maximum number of codewords with pairwise angle above a threshold θ grows as $e^{\alpha(\theta)d}$. The constant α depends on θ ; for the noise tolerance built into modern Hopfield’s energy landscape, $\alpha = 1/4$.

If patterns were chosen by a designed codebook (Reed–Solomon, BCH, lattice codes) rather than placed randomly, α could be pushed closer to $\ln 2 \approx 0.69$. The hierarchy:

Construction	Capacity exponent	Bits per dim
Classical Hopfield (quadratic energy)	$\sim (\log d)/d$	$\rightarrow 0$
Modern Hopfield, random patterns	$\sim d/4$	0.36
Modern Hopfield, designed codebook	$\rightarrow d \ln 2$	$\rightarrow 1.0$
Information-theoretic ceiling	$d \ln 2$	1.0

Transformers sit in the middle row. The key vectors are “placed randomly” in the sense that they emerge from learned projections without any structured codebook design. So attention extracts about 36% of the information-theoretic ceiling on memory per head dimension — for free, with no engineering effort.

Nobody has tried building transformers with designed codebooks for keys. The marginal capacity improvement would be a factor of ~ 3 in bits-per-dimension, which is small compared to what other things (compute, training data, model quality) dominate. Probably not worth the engineering complexity.

Why this matters for “scaling head dim”. The linear cost in d buys exponential memory in d . Doubling d from 64 to 128 does not double capacity — it *squares* it ($10^7 \rightarrow 10^{14}$). That asymmetry is one of the deep architectural wins of attention. The hardware cost of wider heads is linear; the representational capacity gained is exponential. It is the same reason binary trees with d levels can index 2^d leaves — but applied to a continuous geometric construction (the unit sphere).

Convergence of perspectives. Physics gives us the energy function and the partition function. Information theory gives us the Shannon ceiling. Coding theory gives us the random-coding bound and the gap to the ceiling. Modern Hopfield is the same construction read through all three lenses, and they agree on the same exponential scaling with the same coefficient. The 2024 Nobel honoured work sitting at exactly the point where physics, information theory, and machine learning meet — and the math is identical in all three.

10 References

- J. J. Hopfield. *Neural networks and physical systems with emergent collective computational abilities*. PNAS, 1982.
- D. J. Amit, H. Gutfreund, H. Sompolinsky. *Storing infinite numbers of patterns in a spin-glass model of neural networks*. Physical Review Letters, 1985.
- D. Krotov, J. J. Hopfield. *Dense Associative Memory for Pattern Recognition*. NIPS, 2016.
- M. Demircigil, J. Heusel, M. Löwe, S. Upgang, F. Vermet. *On a model of associative memory with huge storage capacity*. Journal of Statistical Physics, 2017.
- H. Ramsauer et al. *Hopfield Networks Is All You Need*. arXiv:2008.02217, 2020.